

Machine Learning Fairness

Group Member:

Gexin Chen

Noah Love

Rhea Sablani

Jiachen Liu

Xubo Wang

Yixuan Zhang

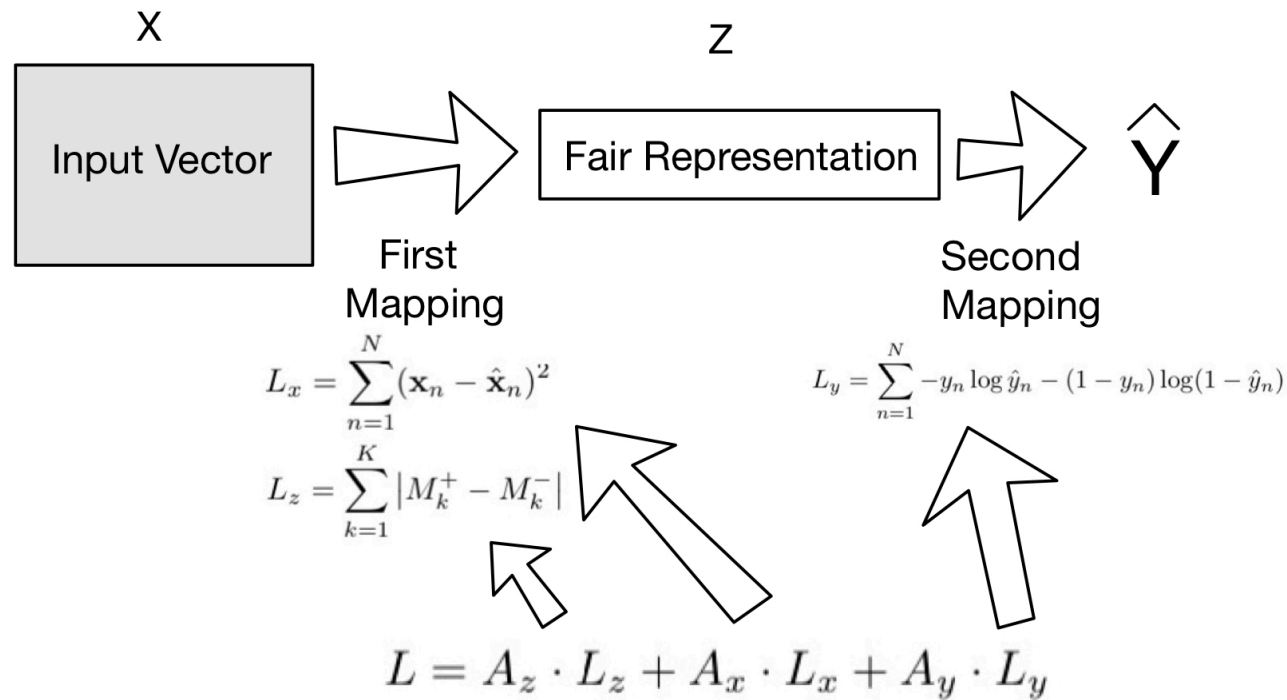
Evaluation Metrics

- **Accuracy** (Proportion of Correct Prediction)
- **Individual Fairness** (Similar Individuals Receive Similar Classification)
- **Group Fairness** (Similar or Equivalent Classification Distribution)
- **Calibration** (Difference of Accuracy Between Groups)

Algorithms

1. **Learning Fair Representations**
2. **Maximizing Accuracy under Fairness Constraints (C-SVM & C-LR)**
3. **Maximizing Fairness under Accuracy Constraints (gamma & fine-gamma)**

Learning Fair Representation



A here is a hyper-parameter controlling the trade-off between accuracy and fairness.

Maximize Accuracy under Fairness Constraint

C-SVM & C-LR

$$\begin{array}{ll} \text{minimize} & L(\boldsymbol{\theta}) \\ \text{subject to} & \frac{1}{N} \sum_{i=1}^N (\mathbf{z}_i - \bar{\mathbf{z}}) d_{\boldsymbol{\theta}}(\mathbf{x}_i) \leq \mathbf{c}, \\ & \frac{1}{N} \sum_{i=1}^N (\mathbf{z}_i - \bar{\mathbf{z}}) d_{\boldsymbol{\theta}}(\mathbf{x}_i) \geq -\mathbf{c}, \end{array}$$

Letter **c** here is a hyper-parameter controlling the trade-off between accuracy and fairness.

Maximize Fairness under Accuracy Constraint

gamma & fine-gamma

$$\begin{array}{ll} \text{minimize} & |\frac{1}{N} \sum_{i=1}^N (\mathbf{z}_i - \bar{\mathbf{z}}) d_{\boldsymbol{\theta}}(\mathbf{x}_i)| \\ \text{subject to} & L(\boldsymbol{\theta}) \leq (1 + \gamma)L(\boldsymbol{\theta}^*), \end{array}$$

Gamma here is a hyper-parameter controlling the trade-off between accuracy and fairness.

Evaluation

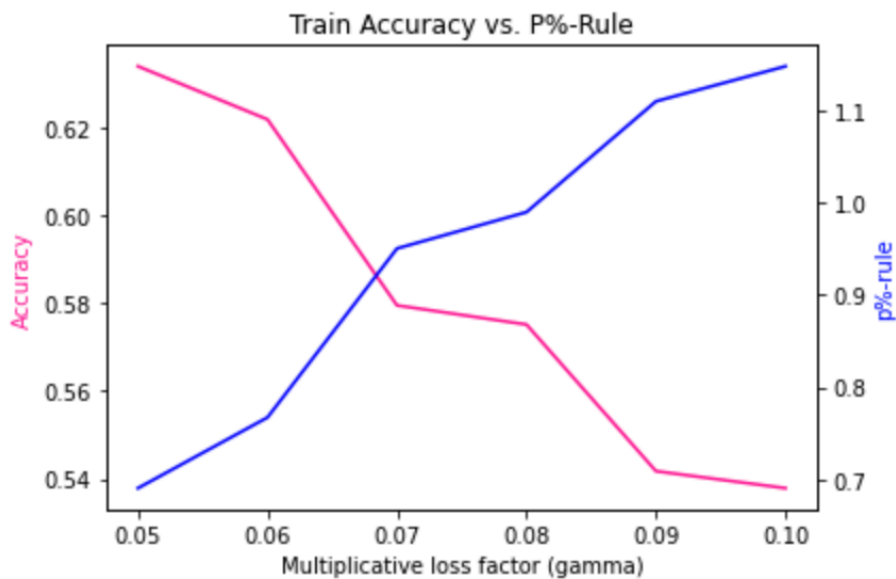
	A1 (LFR)	A2 (C-SVM & C-LR)	A3 (gamma & fine-gamma)
Accuracy	53.66%	66%	56.39%
Calibration	11.77%	4.94%	1.38%
Protected Accuracy	60.69%	63.96%	NA
Non-protected Accuracy	48.91%	68.91%	NA
Individual Fairness	0.8118	NA	NA
p-rule	NA	NA	80%

A1 Under $A_x = 0.001$, $A_y = 0.1$ $A_z = 1000$

A2 Under $c = 1$

A3 Under $\gamma = 0.07$

Hyper-Tuning Visualization



Negative Correlation Between Fairness and Accuracy

Reference

- Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rodriguez, and Krishna P. Gummadi. Fairness Constraints: Mechanisms for Fair Classification. <https://arxiv.org/abs/1507.05259>
- Richard Zemel, Yu Wu, Kevin Swersky, and Toniann Pitassi. Learning Fair Representations. <http://proceedings.mlr.press/v28/zemel13.pdf>